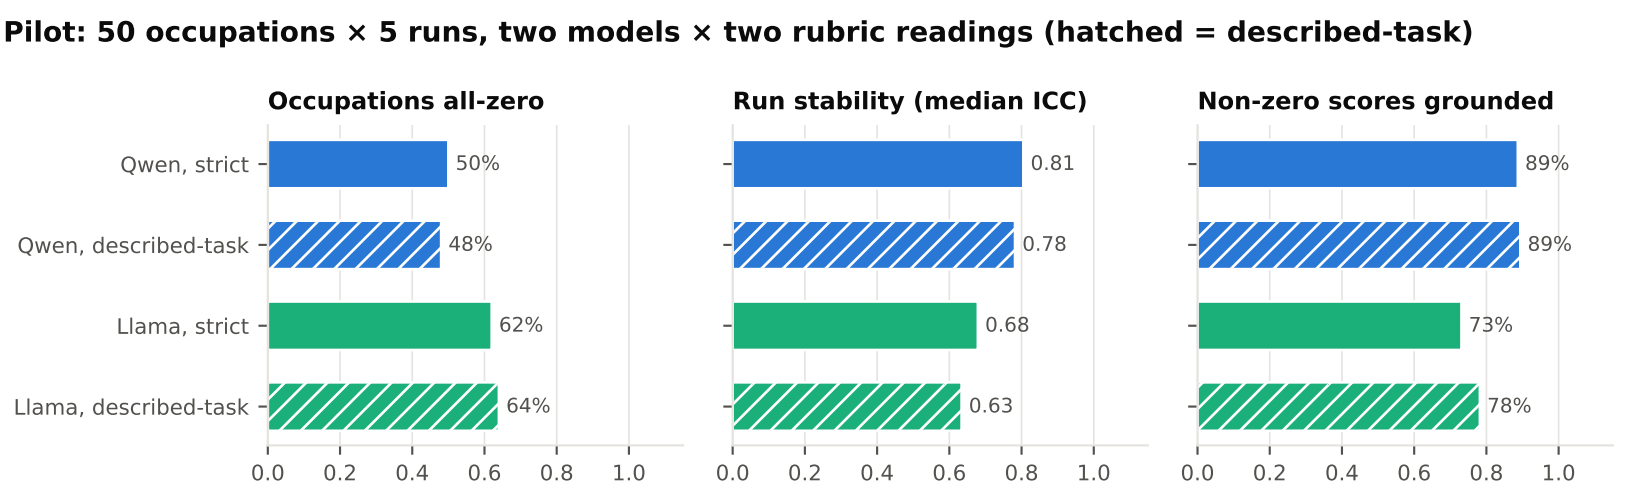
Pilot: 50 occupations × 5 runs, two models × two rubric readings (hatched = described-task)
Occupations all-zero
Run stability (median ICC)
Non-zero scores grounded
Qwen, strict
Qwen, described-task
Llama, strict
Llama, described-task
50%
48%
62%
64%
0.81
0.78
0.68
0.63
89%
89%
73%
78%
0.0
0.2
0.4
0.6
0.8
1.0